

THE FOR GOOD PROJECT

Gap Analysis & Operating Plan

From mission review to a proven, copyable loop

“People and AI, donated to the public good. We research Aotearoa’s hardest problems in the open, and carry the answers through to real impact — in a way any country can copy.”

Adopted mission statement · 2 July 2026

- | | |
|---|---|
| 1 | The verdict — what the mission gets right |
| 2 | The eight gaps |
| 3 | Mission statement — four variations compared |
| 4 | Supply-push vs demand-pull — two operating models |
| 5 | The demand-side plan — founding partner loop |
| 6 | Built to be copied — kernel and instances |
| 7 | How we collaborate — one workspace, two audiences |
| 8 | What we should action next |


Prepared for the working group · thecolab.ai · July 2026 · Companion files: for-good-project-gap-analysis.md, how-we-collaborate.md

1 The verdict — what the mission gets right

The mission is coherent and unusually well-articulated for a project this young. The core mechanism — spare AI capacity meets unworked NZ problems, **agents grind, humans steer**, done means adoption — is sound, and the Constitution / Manifesto / working-docs layering is a genuine strength. Most communities never write down “outcomes over outputs” this crisply.

The gap pattern is equally consistent, and it's one picture:

The gap pattern in one picture



Everything upstream of a merged PR is well-designed; everything between merge and a changed life is aspirational. The Manifesto names the right principles — demand-pull, adoption-as-done, start narrow — but the repo's machinery doesn't yet implement any of them.

The repo proves its own point: ~29 open issues and ~27 open PRs against a handful of humans. Agent output is already outrunning human judgement capacity — exactly the failure mode the Manifesto warns against.

2

The eight gaps

#	Gap	What's missing	Recommendation
1	No demand side	“Demand-pull” is claimed but no mechanism exists for a real organisation to ask a question. Problems look internally invented.	The founding-partner plan (§5). This is the root gap — most others flow from it.
2	Pipeline ends before the mission does	Last stage is Build; “done = adoption” has no stage, owner or instrument. The three declared metrics are tracked nowhere.	Add a Deliver stage + a public IMPACT.md log of adoption events.
3	Breadth contradicts “start narrow”	Five domains, dozens of parallel items, no stream proven end-to-end.	Publicly park all but one stream; drive it through to Deliver, then reopen.
4	Steward supply is the bottleneck	The design rests on scarce human judgement with no recruitment, onboarding or workload story. The PR backlog is this, visible.	Bound the ask (~1 hr/wk); recruit stewards via the partner; cap open agent work per stream.
5	Te Tiriti & Māori data sovereignty absent	First-listed domain is child welfare, yet no founding doc mentions Te Tiriti, iwi partnership or Māori data sovereignty. A legitimacy risk funders will notice immediately.	Constitutional amendment; treat “we need Māori partnership we don't have” as a valid always-human finding.
6	No maintenance story	If an NFP relies on a tool built here, nobody owns its upkeep. Abandonment does harm — violating Article III.	No G2 pass without a named maintenance plan: handover, sunset date, or maintainer.
7	No ecosystem positioning	No stated relationship to Figure.NZ, Hui E!, community law, existing civic-tech. Risk of duplicating; those orgs are also the fastest route to a partner.	A short ECOSYSTEM.md: who already works each domain, where we complement.
8	No definition of “proven” for replication	“Once proven out” is as unfalsifiable as “done” was before it meant adoption.	Explicit replication-readiness criteria (§6).

Gaps 5 and 6 need their own sessions — the first-partner choice sidesteps them temporarily but doesn't resolve them. Both must land before any stream enters child welfare or ships a tool an organisation depends on. The replication ambition raises the bar on both: they'll be closed as kernel patterns other countries inherit, so they're worth authoring carefully, once.

3

Mission statement — four variations compared

The test applied to every candidate: does it carry the four load-bearing ideas — **donated** (the spare-capacity model), **open** (the commons), an arc to **real impact** (done = adoption, not merge), and **copyable** (the replication ambition)?

Variation	Donated	Open	Impact arc	Copyable	Sayable
V1 — Original README “An open research commons where people and AI agents work together to solve NZ's biggest societal problems.”	✗	✓	✗	✗	✓
V2 — Pipeline version “We unite and donate our human expertise and AI capacity to identify society's greatest challenges... lasting impact for the greater good of NZ and beyond.”	✓	✗	✓	~	✗
V3 — Demand-pull version “We donate our expertise and AI capacity to an open commons: taking real problems from NZ communities... sharing the model with the world.”	✓	✓	✓	✓	~
V4 — Adopted ★ “People and AI, donated to the public good. We research Aotearoa's hardest problems in the open, and carry the answers through to real impact — in a way any country can copy.”	✓	✓	✓	✓	✓

Variation	Positives	Negatives
V1 — the original	“Open research commons” is precise and distinctive — it names the mechanism. Short enough to repeat.	Ends at “solve problems”: activity, not outcome — the exact gap the Manifesto warns about. No donation idea, no replication.
V2 — the pipeline version	“Donate” is the standout word nobody else uses. The arc mirrors the full pipeline to impact; “and beyond” carries replication.	Forty words, five clauses. “Greatest challenges” is supply-push framing. Drops “open”. “Greater good / lasting impact” are the clichés every foundation uses.
V3 — the demand-pull version	Only variation to hard-code demand-pull and adoption into the words themselves — mission and operating model fully agree.	Quietly forecloses the supply-push exception: public-interest domains like AI policy have no community asker. Reads as a sentence, not a saying.
V4 — adopted	Carries all four ideas. The opening (“People and AI, donated to the public good”) works standalone as a tagline. “Carry the answers through” is the Deliver stage in plain words.	“Hardest problems” retains a supply-push flavour — acceptable if the asker-or-exploratory rule (§4) is adopted, since the mechanism then enforces what the poetry doesn't.

4

Supply-push vs demand-pull — two operating models

The language question (“society’s greatest challenges” vs “problems real organisations bring us”) isn’t cosmetic — each phrasing implies a different operating model, and each model builds different things. Choose consciously rather than drift.

A — SUPPLY-PUSH

“We identify society’s greatest challenges” · what the repo is today, in practice

What gets built: Broad landscape research, systemic analyses, tools for a *class* of user designed before any specific org is at the table. The domain taxonomy and fan-out machinery are this model’s infrastructure.

Merits: Works from a cold start. Free to tackle important-but-unasked problems (AI policy, civic transparency genuinely have no “customer”). Produces wide-reach public goods. Contributors can start immediately.

Negatives: Nobody is waiting — adoption must be *sold* after the fact, the hardest way. “Greatest” is judged from the inside. Unfalsifiable until a user is found. Burns scarce human judgement on work with no committed beneficiary. The 27-PR backlog is this model’s signature.

B — DEMAND-PULL

“We answer questions real organisations bring us” · what the Manifesto claims

What gets built: Narrow, specific artifacts with a named recipient — “the answer to *this* org’s question”, generalised later. Investment shifts from fan-out machinery to intake, relationship and feedback machinery.

Merits: Adoption structurally built in — the asker is waiting, so “done = adoption” has an obvious test. Relevance validated before tokens are spent. Impact stories write themselves. Naturally rate-limits agent output to human capacity.

Negatives: Cold-start dependent — no partner, no work; partner acquisition is slow human relationship-building. Scope bounded by whoever shows up; systemic problems with no natural asker get missed. Risks becoming free consultancy for one org.

Recommendation — demand-anchored, with a bounded exploratory allowance.


Every stream must **name its asker, or explicitly declare itself exploratory** — and exploratory streams are capped (one active at a time, time-boxed, must convert to a named asker by G1 or park). Public-interest work stays possible; “nobody asked for this” becomes a visible, budgeted exception instead of the invisible default. This rule travels to any country unchanged — it’s a kernel pattern.

5

The demand-side plan — founding partner loop

Decisions taken: first partner is **one small NFP**; intake is a **named human broker** plus a **simple public form**, formalised by a lightweight **partner charter**; starting position is loose contacts requiring deliberate outreach.

The founding-partner loop — the first thing to prove end-to-end



The broker	A named person (a maintainer, initially) owns the relationship: finds and signs the partner, translates their pain into a well-framed Discover issue (the partner never touches GitHub), carries syntheses back in plain language, files their feedback. Cannot be an agent, form or bot at founding stage.
Finding the partner	Run as the project's own method: a Discover issue — “which NZ small NFPs have a research-shaped problem we could answer in 4–6 weeks?” Agents map the candidate landscape; humans work the loose contacts. Grant-access domain first; deliberately not child welfare until the Te Tiriti gap is closed.
The charter	One page, plain language, not a legal contract. Partner gives: a real question, one feedback call per synthesis, honesty when we miss. Project gives: rigorous cited research free, no data misuse, nothing published about them without sign-off, freedom to walk away. Public template in docs/ — makes demand-pull auditable.
The public form	“Ask us a question” on the dashboard → needs-triage Discover candidate. G0 already handles the noise. Secondary channel — expect the founding partner to come from outreach, not the form.

One quarter to a falsifiable north star

success = first IMPACT.md entry, even if it says “partner said we missed”


NOW	WEEKS 1-4	WEEKS 4-8	WEEKS 8-13
· Name broker · open partner issue · draft charter	· Work loose contacts · map candidates · park other streams	· Sign charter · frame partner stream · add IMPACT.md	· Stream through G1 with partner · first impact entry

6

Built to be copied — kernel and instances

The intent to package the model for peers abroad **sharpens the analysis rather than reshaping it**. The exportable asset today would be scaffolding without evidence it produces adoption — exporting that exports the failure mode at scale. The first proven stream *is* the replication prerequisite. Resist building the franchise kit before the first store works.

Close every gap as a kernel pattern, not an NZ patch



The legitimacy layer is a design slot, not an NZ footnote. Every country instance needs one — indigenous data sovereignty, marginalised-community partnership, whatever legitimacy demands there. The kernel makes it a required slot: no sensitive domain opens in any instance until it's documented. NZ's Te Tiriti answer becomes the exemplar others copy — which strengthens the case for doing it properly.

“Proven” means, at minimum:

① one stream through Deliver end-to-end · ② IMPACT.md records genuine adoption events with partner feedback · ③ a written playbook a stranger could follow · ④ the kernel/instance boundary documented · ⑤ the legitimacy slot has a worked NZ example.


Until then, “packaging the model” means writing artifacts as templates as you go — not a separate workstream.

7

How we collaborate — one workspace, two audiences

The principle: GitHub is the single source of truth, and non-technical people never have to see it. They read through rendered views and write through low-friction channels; automation and named humans bridge both directions. The moment content is mirrored into a second tool “for the non-technical folks”, there are two sources of truth and one of them is lying.

One workspace, two audiences



Who	Reads	Writes via
Partner organisation	Their stream's permalink page	The broker (calls); later the form
Frontline worker / SME	Stream overviews on the site	Steward or broker · WhatsApp bot · issue comments if willing
Stream steward	Site + the repo	Direct: edits overviews, makes gate decisions
Contributor (human or agent)	The repo	Direct: issues, branches, PRs
Maintainer	Everything	Direct, plus converting tier 1–3 input into repo mechanics

Sequencing — don't build the bridge before there's traffic. The bot and forms are themselves supply-push infrastructure, tempting to build early. Tier 1 (broker as bridge) works from day one with zero code and is all the founding-partner phase needs. Build the form alongside the demand side's public launch; build the bot when feedback volume actually strains the humans. Sensitive material (partner contacts, unsigned drafts) lives in one small private repo — one rule: *if it isn't cleared for public, it lives there.*

The portable pattern: **Read = rendered site. Write = named humans, then chat, then forms. GitHub = the database. Automation = the glue.** A country instance swaps WhatsApp for whatever channel its community lives in; everything else travels as-is.

8

What we should action next

Immediate targets for the group. Each is small enough to start this week; most are PRs against the repo — arguable in the open, per the Constitution.

#	Action	Detail	When	Type
1	Ratify the mission statement	PR the adopted statement (V4) into README + Manifesto, linking this analysis as the reasoning. Settles the framing everything else hangs off.	This week	PR
2	Adopt the asker-or-exploratory rule	PR to STREAMS.md: every stream names its asker or declares itself exploratory; exploratory capped at one, time-boxed, converts by G1 or parks. Turns demand-pull from prose into mechanism.	This week	PR
3	Name the broker; open the founding-partner issue	One person owns finding + signing the first small NFP. Agents map candidates; humans work the loose contacts. Grant-access first; not child welfare.	This week	People
4	Draft the partner charter template	One page, plain language, in docs/ as a kernel template — needed before any partner conversation gets serious.	This week	PR
5	Publicly park all streams but one	Pick the stream most likely to reach a real user; park the rest citing “depth before breadth”. Reversible, honest, relieves the steward bottleneck.	This month	Decision
6	Triage the open-PR backlog	~27 open PRs is the project's own named failure mode, live in its own repo. Merge, close or park each; then cap open agent work per stream.	This month	Housekeeping
7	Add the Deliver stage + IMPACT.md	PR the fifth pipeline stage and an empty impact log with its entry format. An empty, honest impact log is itself a statement of what we measure.	This month	PR
8	Open the Te Tiriti workstream	State the gap plainly; begin identifying Māori advisors or partner orgs to co-author the legitimacy layer. Must land before child welfare proceeds.	Start now — long pole	People

Items 1–4 are a single week's work for a small group. Items 5–7 are housekeeping with outsized signalling value. Item 8 is the long pole — start it now so it's ready when it's needed.

The whole game: one real person, at one real organisation, measurably better off — then again, and again, until the way we did it is worth copying.